

---

---

# Predicting Diabetes

With Machine Learning

Group 3:

— Luke, Nikhil, Aleksey, Maxiel, Darrick, Samantha, —  
Ruqayyah

---

---

# Introduction

Diabetes is a serious chronic disease in which individuals lose the ability to effectively regulate levels of glucose in the blood and can lead to reduced quality of life and life expectancy.

The Behavioral Risk Factor Surveillance System (BRFSS) is a health-related telephone survey that is collected annually by the CDC.

The Survey collects responses from over 400,000 Americans on health-related risk behaviors, chronic health conditions, and the use of preventative services.

# diabetes\_binary\_5050split\_health\_indicators\_BRFSS2015.csv

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V
	Diabetes	HighBP	HighChol	CholCheck	BMI	Smoker	Stroke	HeartDise	PhysActiv	Fruits	Veggies	HvyAlcoh	AnyHealth	NoDocbc	GenHlth	MentHlth	PhysHlth	DiffWalk	Sex	Age	Education	Income
	0	1	0	1	26	0	0	0	1	0	1	0	1	0	3	5	30	0	1	4	6	8
	0	1	1	1	26	1	1	0	0	1	0	0	1	0	3	0	0	0	1	12	6	8
	0	0	0	1	26	0	0	0	1	1	1	0	1	0	1	0	10	0	1	13	6	8
	0	1	1	1	28	1	0	0	1	1	1	0	1	0	3	0	3	0	1	11	6	8
	0	0	0	1	29	1	0	0	1	1	1	0	1	0	2	0	0	0	0	8	5	8
	0	0	0	1	18	0	0	0	1	1	1	0	0	0	2	7	0	0	0	1	4	7
	0	0	1	1	26	1	0	0	1	1	1	1	1	0	1	0	0	0	1	13	5	6
	0	0	0	1	31	1	0	0	0	1	1	0	1	0	4	0	0	0	1	6	4	3
	0	0	0	1	32	0	0	0	1	1	1	0	1	0	3	0	0	0	0	3	6	8
	0	0	0	1	27	1	0	0	0	1	1	0	1	0	3	0	6	0	1	6	4	4
	0	1	1	1	24	1	0	1	1	1	1	0	1	0	3	0	4	0	0	12	4	6

# diabetes\_binary\_health\_indicators\_BRFSS2015.csv

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X
	Diabetes_binary	HighBP	HighChol	CholCheck	BMI	Smoker	Stroke	HeartDisease	PhysActivity	Fruits	Veggies	HvyAlcoholC	AnyHealth	NoDocbc	GenHlth	MentHlth	PhysHlth	DiffWalk	Sex	Age	Education	Income		
1	0	1	0	1	40	1	0	0	0	0	1	0	1	0	5	18	15	1	0	9	4	3		
2	0	1	0	0	25	1	0	0	1	0	0	0	0	1	3	0	0	0	0	7	6	1		
3	0	1	1	1	28	0	0	0	0	1	0	0	1	1	5	30	30	1	0	9	4	8		
4	0	1	0	1	27	0	0	0	1	1	1	0	1	0	2	0	0	0	0	11	3	6		
5	0	1	1	1	24	0	0	0	1	1	1	0	1	0	2	3	0	0	0	11	5	4		
6	0	1	1	1	25	1	0	0	1	1	1	0	1	0	2	0	2	0	1	10	6	8		
7	0	1	0	1	30	1	0	0	0	0	0	0	1	0	3	0	14	0	0	9	6	7		
8	0	1	1	1	25	1	0	0	1	0	1	0	1	0	3	0	0	1	0	11	4	4		
9	1	1	1	1	30	1	0	1	0	1	1	0	1	0	5	30	30	1	0	9	5	1		
10	0	0	0	1	24	0	0	0	0	0	1	0	1	0	2	0	0	0	1	8	4	3		
11	1	0	0	1	25	1	0	0	1	1	1	0	1	0	3	0	0	0	1	13	6	8		
12	0	1	1	1	34	1	0	0	0	1	1	0	1	0	3	0	30	1	0	10	5	1		
13	0	0	0	1	26	1	0	0	0	0	1	0	1	0	3	0	15	0	0	7	5	7		
14	1	1	1	1	28	0	0	0	0	0	1	0	1	0	4	0	0	1	0	11	4	6		
15	0	0	1	1	33	1	1	0	1	0	1	0	1	1	4	30	28	0	0	4	6	2		
16	0	1	0	1	33	0	0	0	1	0	0	0	1	0	2	5	0	0	0	6	6	8		
17	0	1	1	1	21	0	0	0	1	1	1	0	1	0	3	0	0	0	0	10	4	3		
18	1	0	0	1	23	1	0	0	1	0	0	0	1	0	2	0	0	0	1	7	5	6		

# Supervised Learning Model

We attempted to analyze the data  
with a supervised model

---

- We scaled the non-binary variables to ensure that all features contribute equally to the model training process and to prevent certain variables from dominating the learning process due to differences in their scales

```
In [8]: df_data_scaled = StandardScaler().fit_transform(diabetes[['BMI', 'GenHlth',  
    'MentHlth', 'PhysHlth', 'Age', 'Education',  
    'Income']])
```

```
In [9]: df_data_scaled = pd.DataFrame(df_data_scaled, columns=['BMI', 'GenHlth',  
    'MentHlth', 'PhysHlth', 'Age', 'Education',  
    'Income'])
```

- We then created the labels, setting (`y``) from the “Diabetes\_binary” column, and then created the features (`X``) DataFrame from the remaining columns.
- We ran the model many times, dropping variables that did not add accuracy to the model

```
0]: # Separate the data into labels and features
```

```
# Separate the y variable, the labels
```

```
y = diabetes["Diabetes_binary"]
```

```
# Separate the X variable, the features
```

```
X = diabetes.drop(columns=["Diabetes_binary", 'GenHlth',  
                           'MentHlth', 'PhysHlth', 'Sex', 'Fruits', 'Veggies', 'Smoker'])
```

```
1]: X = pd.concat([X, df_data_scaled], axis=1)
```

```
2]: # Review the y variable Series
```

```
y.head()
```

```
2]: 0    0.0
```

```
1    0.0
```

```
2    0.0
```

```
3    0.0
```

```
4    0.0
```

```
Name: Diabetes_binary, dtype: float64
```

- The balanced 50/50 data set was used, so there was no further work necessary to balance the data
- We then split the data into a training and test split and ran the logistic regression model
- A balanced accuracy score of 0.749 means that, on average, the model correctly predicts the target variable for about 74.9% of the instances

```
# Print the balanced_accuracy score of the model  
balanced_accuracy_score(y_test, testing_predictions)
```

```
0.7492281552726585
```

## Logistic Regression with the Original Data

- Precision: Model is correct about 76% of the time for class 0.0 and 74% for class 1.0.
- Recall: Correctly identifies about 74% of instances for class 0.0 and 76% for class 1.0.
- F1-Score: Balanced performance with F1-scores around 0.75 for both classes.
- Overall Accuracy: Model correctly predicts target variable for 75% of instances.

```
In [20]: # Print the classification report for the model
testing_report = classification_report(y_test, testing_predictions)
print(testing_report)
```

	precision	recall	f1-score	support
0.0	0.76	0.74	0.75	8913
1.0	0.74	0.76	0.75	8760
accuracy			0.75	17673
macro avg	0.75	0.75	0.75	17673
weighted avg	0.75	0.75	0.75	17673



# Neural Network Learning Model

---

Original NN  
Built W/O  
Optimizer

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====		
dense_4 (Dense)	(None, 64)	1408
dense_5 (Dense)	(None, 32)	2080
dense_6 (Dense)	(None, 16)	528
dense_7 (Dense)	(None, 1)	17
=====		
Total params: 4033 (15.75 KB)		
Trainable params: 4033 (15.75 KB)		
Non-trainable params: 0 (0.00 Byte)		
=====		

553/553 - 1s - loss: 0.5120 - accuracy: 0.7479 - 783ms/epoch - 1ms/step  
Loss: 0.511986494064331, Accuracy: 0.7478639483451843

```
top_model = tuner.get_best_models(3)
for model in top_model:
    model_loss, model_accuracy = model.evaluate(X_test_scaled,y_test,verbose=2)
    print(f"Loss: {model_loss}, Accuracy: {model_accuracy}")
```

```
553/553 - 1s - loss: 0.5055 - accuracy: 0.7541 - 804ms/epoch - 1ms/step
Loss: 0.5055218935012817, Accuracy: 0.7541447281837463
553/553 - 1s - loss: 0.5019 - accuracy: 0.7540 - 744ms/epoch - 1ms/step
Loss: 0.5019425749778748, Accuracy: 0.753974974155426
553/553 - 1s - loss: 0.5040 - accuracy: 0.7538 - 870ms/epoch - 2ms/step
Loss: 0.503968358039856, Accuracy: 0.7538052201271057
```

Optimizer Data, Top three models  
and best model values

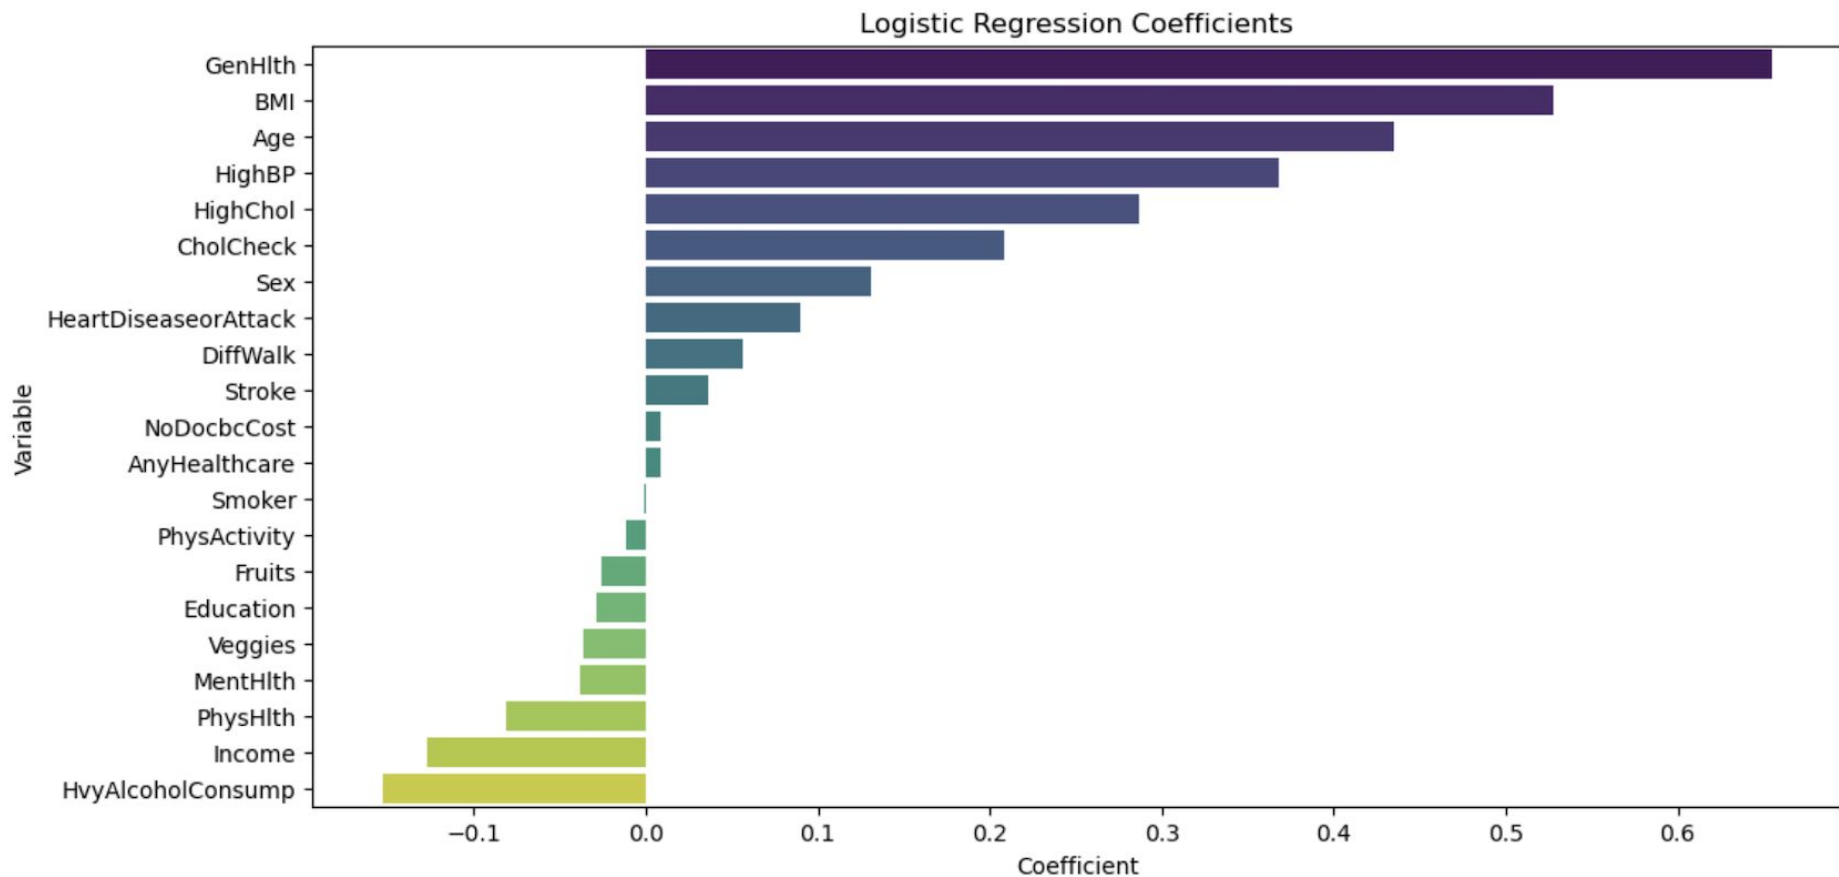
best\_hyper.values

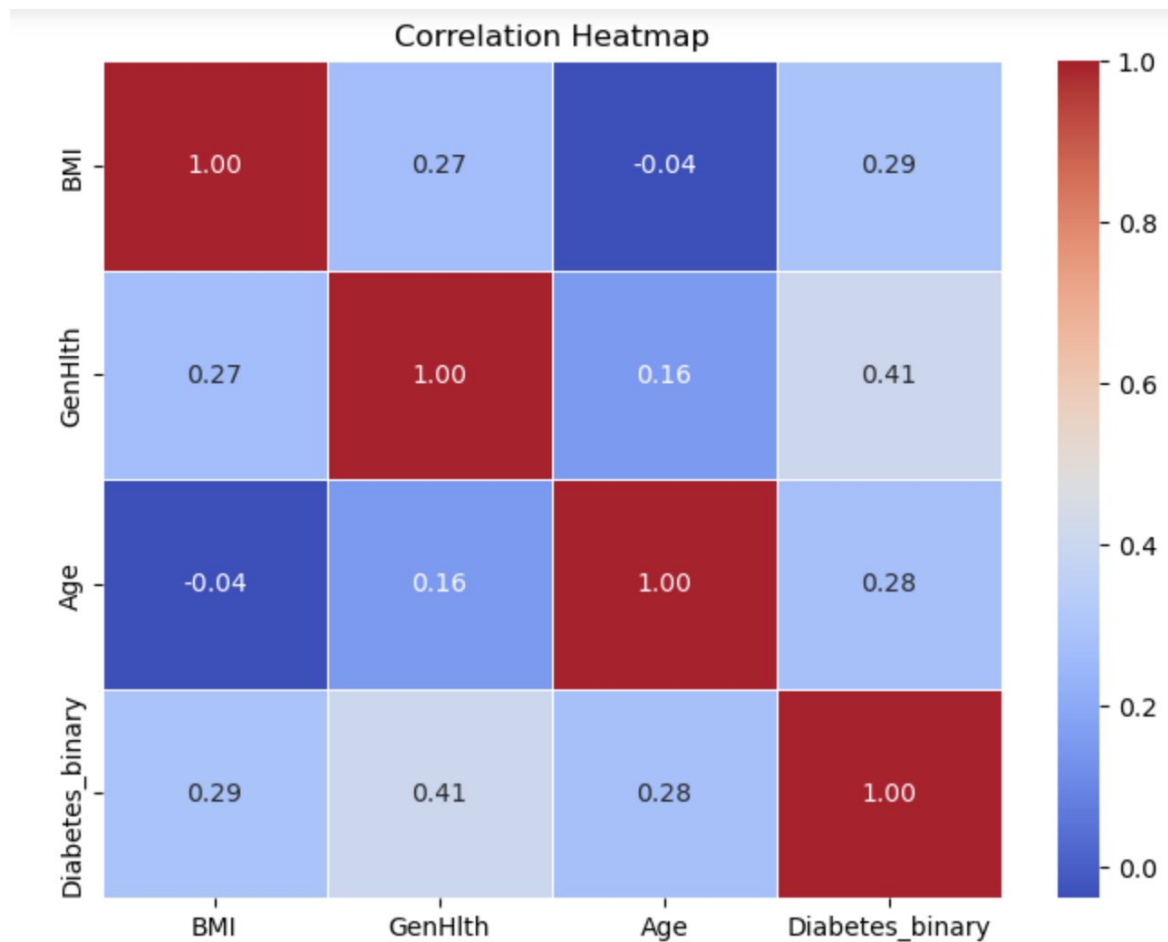
```
{'activation': 'relu',
 'first_units': 9,
 'num_layers': 4,
 'units0': 9,
 'units1': 1,
 'units2': 3,
 'units3': 9,
 'units4': 1,
 'units5': 1,
 'tuner/epochs': 7,
 'tuner/initial_epoch': 3,
 'tuner/bracket': 2,
 'tuner/round': 1,
 'tuner/trial_id': '0036'}
```

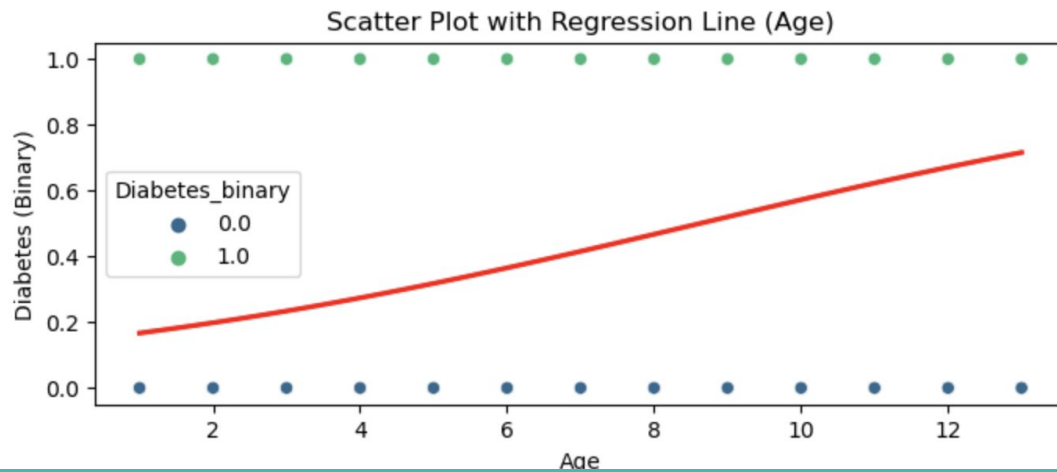


553/553 - 1s - loss: 0.5041 - accuracy: 0.7530 - 1s/epoch - 2ms/step  
Loss: 0.5041230916976929, Accuracy: 0.7529565095901489

	precision	recall	f1-score	support	Model: "sequential"		
0.0	0.76	0.73	0.74	8835	Layer (type)	Output Shape	Param #
1.0	0.74	0.77	0.75	8838	=====		
accuracy			0.75	17673	dense (Dense)	(None, 9)	198
macro avg	0.75	0.75	0.75	17673	dense_1 (Dense)	(None, 1)	10
weighted avg	0.75	0.75	0.75	17673	dense_2 (Dense)	(None, 3)	6
					dense_3 (Dense)	(None, 9)	36
					dense_4 (Dense)	(None, 1)	10
					=====		
					Total params: 260 (1.02 KB)		
					Trainable params: 260 (1.02 KB)		
					Non-trainable params: 0 (0.00 Byte)		


# Diabetes Risk Factors









BMI vs Healthcare Coverage




BMI Group % of Healthcare Coverage



BMI Group % of Heavy Alcohol Consumption



BMI Group Average Mental Health Score





## ***Limitations & Challenges***

- ❑ Imbalanced Dataset.  
Originally had accuracy of about 80%.
- ❑ The balanced dataset had an initial accuracy of 70%

*Thank you*

